

Parallel Graph Algorithms

Design and Analysis of Parallel Algorithms

5DV050 Spring 2012

Part I

Introduction

Overview

- ▶ Graphs—definitions, properties, representation
- ▶ Minimal spanning tree
 - ▶ Prim's algorithm
- ▶ Shortest paths (1-to-all)
 - ▶ Dijkstra's algorithm
- ▶ Shortest paths (all-to-all)
 - ▶ Algorithm based on matrix multiplication
 - ▶ Dijkstra's algorithm
 - ▶ Source partitioned
 - ▶ Source parallel
 - ▶ Floyd's algorithm
- ▶ Transitive closure
- ▶ Connected components

Graphs: Definitions

Graphs

- ▶ $G = (V, E)$: V is the set of *vertices* and E is the set of *edges*

Undirected graphs


- ▶ An *edge* $e \in E$ is an unordered pair $\{u, v\}$ where $u, v \in V$

Directed graphs


- ▶ An *arc* $e \in E$ is an ordered pair (u, v) directed from u to v
- ▶ A *path* from u to v is a sequence u, \dots, v of vertices where consecutive vertices correspond to an arc
 - ▶ *Simple path*: all vertices are distinct
 - ▶ *Cycle*: $u = v$
 - ▶ *Acyclic*: contains no cycles

Examples of graphs

Undirected



Directed



Graphs: Properties

- ▶ A graph is *connected* if it exists a path between every pair of vertices
- ▶ A graph is *complete* if it exists an edge between every pair of vertices
- ▶ $G' = (V', E')$ is a *sub-graph* of $G = (V, E)$ if $V' \subseteq V$ and $E' \subseteq E$
- ▶ A *tree* is a connected acyclic graph
- ▶ A *forest* consists of several trees
- ▶ A graph $G = (V, E)$ is *sparse* if $|E|$ is much smaller than $|V|^2$


Weighted graphs:

- ▶ $G = (V, E, w)$, where w is a real-valued function defined on E (every edge/arc has a value)
- ▶ The weigh of a graph is the sum of the weights of its edges

Matrix representation of graphs

Non-weighted graphs:

$$a_{i,j} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E, \\ 0 & \text{otherwise.} \end{cases}$$



Weighted graphs:



$$a_{i,j} = \begin{cases} w(v_i, v_j) & \text{if } (v_i, v_j) \in E, \\ 0 & \text{if } i = j, \\ \infty & \text{otherwise.} \end{cases}$$

$$A = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

Suitable for dense graphs

List representation of graphs

- ▶ $G = (V, E)$ is represented by the list $\text{Adj}[1 \dots |V|]$ of lists
- ▶ For each $v \in V$, $\text{Adj}[v]$ is a linked list of all vertices that has an edge in common with v



Suitable for sparse graphs

Part II

Minimum Spanning Tree (MST)

Minimum Spanning Tree (MST)


- ▶ A *spanning tree* of a graph $G = (V, E)$ is a sub-graph of G that is a tree and contains all vertices of G
- ▶ MST for a weighted graph is a spanning tree with minimum weight
- ▶ If $G = (V, E)$ is not connected, then it cannot have an MST but instead has a minimum spanning forest
- ▶ Assume that G is connected, otherwise we find connected components and find an MST Of each component

Prim's algorithm

- ▶ Greedy algorithm: Add one edge at a time to the MST
- ▶ Select a vertex at random
- ▶ Choose an edge with minimum weight between the selected set and the unselected set (break ties arbitrarily)
- ▶ Select the unselected vertex
- ▶ Repeat until a spanning tree has been created
- ▶ The constructed tree is an MST


Prim's algorithm: Example (1/6)

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>d</i> []	1	0	5	1	∞	∞




Prim's algorithm: Example (2/6)

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>d</i> []	1	0	2	1	4	∞




Prim's algorithm: Example (3/6)

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>d</i> []	1	0	2	1	4	2




Prim's algorithm: Example (4/6)

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>d</i> []	1	0	2	1	1	2




Prim's algorithm: Example (5/6)

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>d</i> []	1	0	2	1	1	2



Prim's algorithm: Example (6/6)

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>
<i>d</i> []	1	0	2	1	1	2



Prim's algorithm

```
1: PRIM( $V, E, w, r$ )
2:  $V_T := \{r\}$ 
3:  $d[r] := 0$ 
4: for all  $v \in (V - V_T)$  do
5:    $d[v] := w(r, v)$  if  $(r, v) \in E$  else  $d[v] := \infty$ 
6: end for
7: while  $V_T \neq V$  do
8:   Find vertex  $u \in (V - V_T)$  such that
      $d[u] = \min\{d[v] \mid v \in (V - V_T)\}$ 
9:    $V_T := V_T \cup \{u\}$ 
10:  for all  $v \in (V - V_T)$  do
11:     $d[v] := \min\{d[v], w(u, v)\}$ 
12:  end for
13: end while
```

Parallelizing Prim's algorithm

- ▶ $d[v]$ is updated for all v
 - ▶ Cannot choose two vertices in parallel
 - ▶ Cannot parallelize outer `while` loop
 - ▶ Instead we parallelize the inner `for` loop
- ▶ Every process holds a block column of adjacency matrix A :

$$A = \begin{bmatrix} A_1 & A_2 & \cdots & A_p \end{bmatrix}$$

and corresponding part of vector d

- ▶ Process P_i holds vertex subset V_i
- ▶ Owner computes: Process P_i responsible for updating its part of d
- ▶ Find global minimum (line 8) with all-reduce
- ▶ Update d in parallel (line 10)

Analysis of the parallel Prim's algorithm

Per iteration:

- ▶ Computation (line 10): $\Theta(n/p)$
- ▶ All-reduce (line 8): $\Theta(\log_2 p)$

Total for n iterations:

- ▶ $T_P(n, p) = \Theta(n^2/p) + \Theta(n \log_2 p)$
- ▶ $T_S(n) = \Theta(n^2)$
- ▶ Iso-efficiency condition: $n = \Omega(p \log_2 p)$

Part III

Shortest paths (1-to-all)

Dijkstra's algorithm

Dijkstra's algorithm:

- ▶ Essentially identical to Prim's algorithm, except
 - ▶ Instead of $d[u]$ store $\ell[u]$, which is the total weight from r to u

Parallel Dijkstra's algorithm:

- ▶ Identical to the parallel Prim's algorithm (with the change above)
- ▶ Analysis identical

Part IV

Shortest paths (all-to-all)

Shortest paths (all-to-all)

- ▶ Goal is to find the weight of the shortest path between all pairs of vertices
- ▶ The result is a square matrix $D = (d_{i,j})$ where $d_{i,j}$ is the weight of the shortest path from v_i to v_j

Algorithm based on matrix multiplication

- ▶ Let $G = (V, E, w)$ be represented by the matrix A
- ▶ Let $d_{i,j}^{(k)}$ represent the weight of the shortest path from v_i to v_j that contains *a maximum of* k edges
- ▶ (Thus, $D_{i,j}^{(1)} = A$)
- ▶ Let v_m be a vertex in that path
- ▶ Then $d_{i,j}^{(k)} = \min_m \{d_{i,m}^{(k-1)} + w(v_m, v_j)\}$

Matrix multiplication-based algorithm (continued)

- ▶ $D^{(k)}$ computed from $D^{(k-1)}$ using modified matrix multiplication:

$$c_{i,j} := \min_k a_{i,k} + b_{k,j}$$

(Find k that minimizes $a_{i,k} + b_{k,j}$)

- ▶ $D^{(k)} = \underbrace{AA \cdots A}_{k \text{ factors}}$

- ▶ But we *only need* $D^{(n-1)}$
- ▶ Compute $D^{(n-1)}$ using repeated squaring:

$$D^{(1)} \mapsto D^{(2)} \mapsto D^{(4)} \mapsto D^{(8)} \mapsto \dots \mapsto D^{(n-1)}$$

- ▶ Complexity for matrix multiplication: $\Theta(n^3)$
- ▶ Number of steps: $\Theta(\log_2 n)$
- ▶ Total complexity: $T_P(n, p) = \Theta(n^3 \log_2 n)$

Dijkstra's algorithm applied to all-to-all shortest paths

Source-partitioned approach:

- ▶ Distribute the sources
- ▶ Run *sequential* 1-vertex Dijkstra on all processors in parallel
- ▶ Complexity: $\Theta(n^2) \cdot \Theta(n/p) = \Theta(n^3/p)$
- ▶ Perfectly parallel
- ▶ Degree of concurrency limited to $p = \mathcal{O}(n)$
- ▶ Each processor must have access to the entire graph

Source-parallel approach:

- ▶ Partition processors into groups (e.g., of size \sqrt{p})
- ▶ Distributed sources over the groups
- ▶ Run *parallel* 1-vertex Dijkstra on all processor *groups* in parallel
- ▶ Complexity: $[\Theta(n^2/\sqrt{p}) + \Theta(n \log_2 \sqrt{p})] n/\sqrt{p} = \Theta(n^3/p) + \Theta((n^2/\sqrt{p}) \log_2 p)$
- ▶ Degree of concurrency $p = \mathcal{O}(n^2)$ (much better)
- ▶ Each processor needs access only to a sub-graph

Floyd's algorithm

- ▶ Given $G = (V, E, w)$
- ▶ Let $V_k := \{v_1, \dots, v_k\}$ (first k vertices of G)
- ▶ For any pair $v_i, v_j \in V$, consider all paths whose intermediate vertices belong to the subset V_k
- ▶ Let $p_{i,j}^{(k)}$ be the shortest such path and let $d_{i,j}^{(k)}$ be the corresponding weight
- ▶ If the vertex v_k is **not** in the path, then $p_{i,j}^{(k)} = p_{i,j}^{(k-1)}$
- ▶ If the vertex v_k **is** in the path, then it can be split into two paths: One from v_i to v_k and one from v_k to v_j where both paths uses vertices only from V_{k-1}
- ▶ In that case, the weight of the path is $d_{i,j}^{(k)} := d_{i,k}^{(k-1)} + d_{k,j}^{(k-1)}$

Floyd's algorithm

```
1: FLOYD(A)
2:  $D^{(0)} := A$ 
3: for  $k = 1$  to  $n$  do
4:   for  $i = 1$  to  $n$  do
5:     for  $j = 1$  to  $n$  do
6:        $d_{i,j}^{(k)} := \min\{d_{i,j}^{(k-1)}, d_{i,k}^{(k-1)} + d_{k,j}^{(k-1)}\}$ 
7:     end for
8:   end for
9: end for
```

- ▶ Similar in structure to matrix multiplication
- ▶ Parallelized using SUMMA-style algorithm

Part V

Transitive closure

Transitive closure

- ▶ If $G = (V, E)$ is a graph, then its *transitive closure* is the graph $G^* = (V, E^*)$, where

$$E^* := \{(v_i, v_j) \mid \text{exists path from } v_i \text{ to } v_j \text{ in } G\}$$

- ▶ Computes the connectivity matrix A^* such that $a_{i,j} = 1$ if $i = j$ or a path from v_i to v_j exists in G and $a_{i,j} = \infty$ otherwise

Method 1:

- ▶ Set the weights in G to 1 and compute all-pairs shortest paths followed by re-interpreting the output

Method 2:

- ▶ Modify Floyd's algorithm by replacing min with *logical or* and + with *logical and*:


$$d_{i,j}^{(k)} := d_{i,j}^{(k-1)} \quad \text{or} \quad (d_{i,k}^{(k-1)} \quad \text{and} \quad d_{k,j}^{(k-1)})$$

Part VI

Connected components

Connected components

- ▶ Partition V into maximal disjoint subsets C_1, C_2, \dots, C_r such that $V = C_1 \cup C_2 \cup \dots \cup C_r$ and $u, v \in C_i$ if and only if u is reachable from v and vice versa
- ▶ The graph below has three connected components



Depth-first search based algorithm

- ▶ Perform depth-first traversal of the graph to generate a spanning forest
- ▶ Each tree in the forest defines a connected component

Parallel formulation:

- ▶ Give sub-graph $G_i = (V, E_i)$ to process P_i
- ▶ Perform sequential algorithm on each sub-graph in parallel
- ▶ Merge forests pair-wise using $\log_2 p$ steps

Forest merging

- ▶ $\text{find}(x)$: Tree to which x belongs
- ▶ $\text{union}(u, v)$: Merge trees to which u and v belong
- ▶ Merge forest A into forest B :
 - ▶ Send A to processor holding B
 - ▶ For each edge (u, v) (there are at most $n - 1$) in A :
 - ▶ $\text{find}(u)$ in B
 - ▶ $\text{find}(v)$ in B
 - ▶ Same tree? Do nothing
 - ▶ Different trees? $\text{union}(u, v)$ in B
 - ▶ Discard A , continue with B
- ▶ Using appropriate set data structure and algorithms, $\text{find}(x)$ and $\text{union}(u, v)$ have expected constant time complexity
- ▶ Complexity (1D block partitioning):

$$T_P(n, p) = \Theta(n^2/p) + \Theta(n \log_2 p)$$

Part VII

Johnson's algorithm

Johnson's algorithm (1-to-all shortest paths)

- ▶ Dijkstra's algorithm
 - ▶ Find unprocessed u such that

$$d[u] = \min\{d[v] \mid v \text{ unprocessed}\}$$

- ▶ Update for all unprocessed vertices v :

$$d[v] := \min\{d[v], d[u] + w(u, v)\}$$

- ▶ For *sparse graph*, store unprocessed vertices in a *priority queue* based on $d[v]$ (smallest first)
- ▶ Take minimum weight vertex from the priority queue
- ▶ Update adjacent vertices

Johnson's algorithm

```
1: JOHNSON( $V, E, r$ )
2:  $Q := V$ 
3: for all  $v \in Q$  do
4:    $d[v] := \infty$ 
5: end for
6:  $d[r] := 0$ 
7: while  $Q \neq \emptyset$  do
8:    $u := \text{ExtractMin}(Q)$ 
9:   for each  $v$  adjacent to  $u$  do
10:    if  $v \in Q$  and  $d[u] + w(u, v) < d[v]$  then
11:       $d[v] := d[u] + w(u, v)$ 
12:    end if
13:  end for
14: end while
```

Parallel Johnson's (centralized queue)

- ▶ Maintain Q at a centralized location
- ▶ Processors compute new values and request updates of Q
- ▶ Major bottleneck
- ▶ No asymptotic speedup, since $\mathcal{O}(|E|)$ updates that are serialized and each take time $\mathcal{O}(\log_2 n)$ leading to the same complexity as the sequential formulation
- ▶ Moreover, only $|E|/|V|$ vertices can be updated in parallel in each iteration, and this number is small since the graph is assumed sparse

Parallel Johnson's (distributed queue)

Distributed queue:

- ▶ Manage Q using distributed algorithm
- ▶ Requires a machine with very low latency to be practical
- ▶ Even if the update complexity is reduced from $\mathcal{O}(\log_2 n)$ to $\mathcal{O}(1)$, we can expect no more than a $\mathcal{O}(\log_2 n)$ speedup since the updates are applied sequentially

Safe vertices:

- ▶ Let u be a vertex with minimal $d[u]$
- ▶ All vertices v with $d[v] = d[u]$ can be processed in parallel
- ▶ If we know that the minimum edge weight is m , then we can relax this to all vertices v with

$$d[v] \leq d[u] + m$$

Parallel Johnson's (unsafe vertices)

- ▶ Process also unsafe vertices in parallel
- ▶ Leads to an algorithm that is no longer equivalent to the sequential algorithm
- ▶ Process p vertices at the top of Q in parallel
- ▶ Each process maintains its own priority queue
- ▶ The distances might no longer correspond to the shortest paths
- ▶ Detect instances of wrongly computed distances and re-process the corresponding vertices


Part VIII

Weighted matchings

Weighted matchings

- ▶ A *matching* $M(G)$ of a graph $G = (V, E)$ is any sub-graph of G where each vertex is incident to at most one edge
- ▶ Let $w(e)$ be the weight of an edge e
- ▶ We define the weight $w(G)$ of a graph G as the sum of all edge weights
- ▶ A *maximum weighted matching* $M^*(G)$ of G is a matching whose weight $w(M^*(G))$ is maximum among all matchings of G

Matchings and parallel graph partitioning



Sequential greedy weighted matching algorithm

```
1: MATCHING( $V, E$ )  
2:  $M(G) := \emptyset$   
3: while  $E \neq \emptyset$  do  
4:   Pick locally heaviest edge  $e$  from  $E$   
5:   Add  $e$  to  $M(G)$   
6:   Remove  $e = \{u, v\}$  and all edges incident to  $u$  and  $v$  from  $E$   
7: end while
```

Approximates a maximal matching within a factor 2

Parallel greedy weighted matching algorithm

For each vertex v in parallel:

- 1: PMATCHING(V, E)
- 2: $R := \emptyset$
- 3: Initialize N to the neighborhood of v (all vertices adjacent to v)
- 4: Let the *candidate* c be the vertex connected to v by the *locally heaviest edge* in N
- 5: **if** $c \neq \perp$ **then**
- 6: Send req to c
- 7: **end if**
- 8: **while** $N \neq \emptyset$ **do**
- 9: Receive message m from vertex u
- 10: If $m = \text{req}$, then $R := R \cup \{u\}$
- 11: If $m = \text{drop}$, then $N := N - \{u\}$ and update c if $u = c$ and if $c \neq \perp$ send req to c
- 12: If $c \neq \perp$ and $c \in R$, send drop to all vertices in N except c
- 13: **end while**